

EAI Lab6

Quantization

Advisor : CCTsai


TA : 林泳陞

Outline

- Introduction
 - Quantization
 - PTQ and QAT
- Task
 - Example code
 - Task

Introduction - Quantization

- Quantization is a method used to reduce the model size and computation requirements by replacing floating-point 32-bit (FP32) representations with lower bit precision, such as FP16, INT8 or even less bits. However, reducing the number of bits used to represent values can lead to a loss of accuracy. The primary objective is to strike a balance between accuracy and computational resources.



Introduction - Quantization

- The simplest way to implement quantization is to map the original values to an integer range of 0 to 255 or -128 to 127, and then map these 256 integers back to the original value range to minimize the error.

For example:

```
FP32: tensor([ 0.6004, -1.0151,  1.1885,  0.9948, -0.3187, -1.1111,  0.7827, -1.1217])
INT8: tensor([ 0.5979, -1.0147,  1.1868,  0.9966, -0.3171, -1.1143,  0.7791, -1.1234])
INT8 integer representation: tensor([ 62., -116., 127., 106., -39., -127., 82., -128.])
```

- Therefore, we need to find a method to implement quantization using scale factor and zero point.

Introduction - Quantization

- $q = \text{round}(\frac{r}{s} + Z)$, where $s = \frac{\beta - \alpha}{\beta_q - \alpha_q}$ and $Z = \text{round}(\alpha_q - \frac{\alpha}{s})$
- For weight, $[\alpha, \beta]$ is weight range, and $[\alpha_q, \beta_q]$ is usually $[-128, 127]$ for int8 quantization.
- For activation, $[\alpha, \beta]$ is output range, and $[\alpha_q, \beta_q]$ is usually $[-128, 127]$ for int8 quantization.
- $r_q = (q - Z) * s$

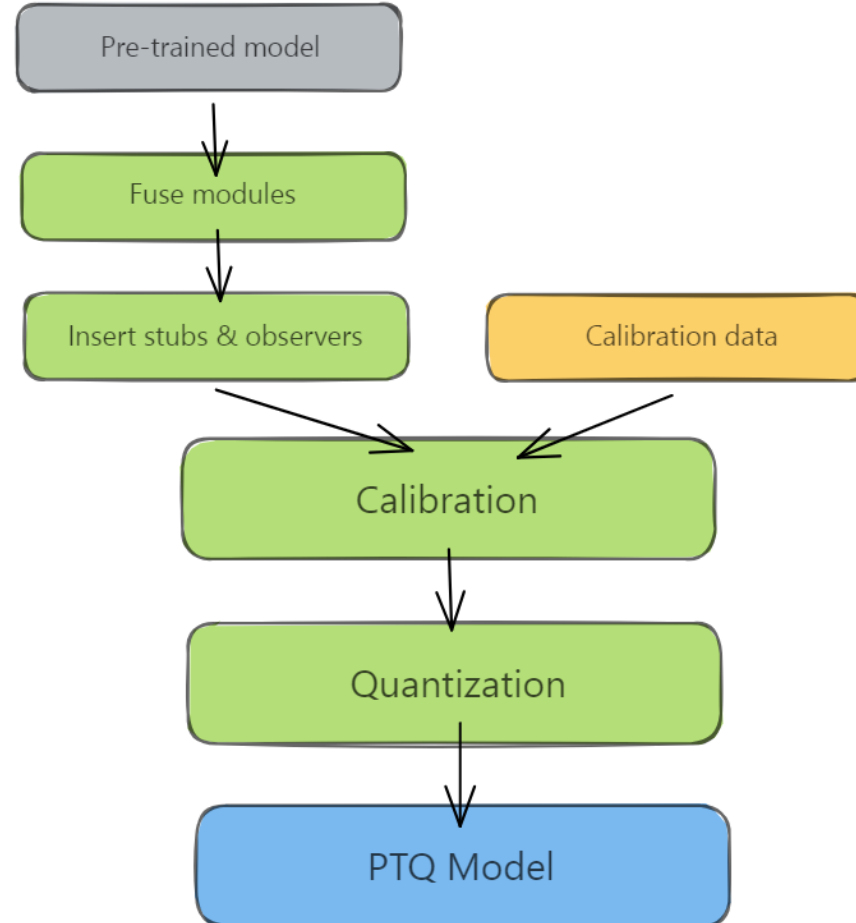
```
FP32: tensor([ 0.6004, -1.0151,  1.1885,  0.9948, -0.3187, -1.1111,  0.7827, -1.1217])
INT8: tensor([ 0.5979, -1.0147,  1.1868,  0.9966, -0.3171, -1.1143,  0.7791, -1.1234])
INT8 integer representation: tensor([ 62., -116., 127., 106., -39., -127.,  82., -128.])
```

Introduction - Quantization

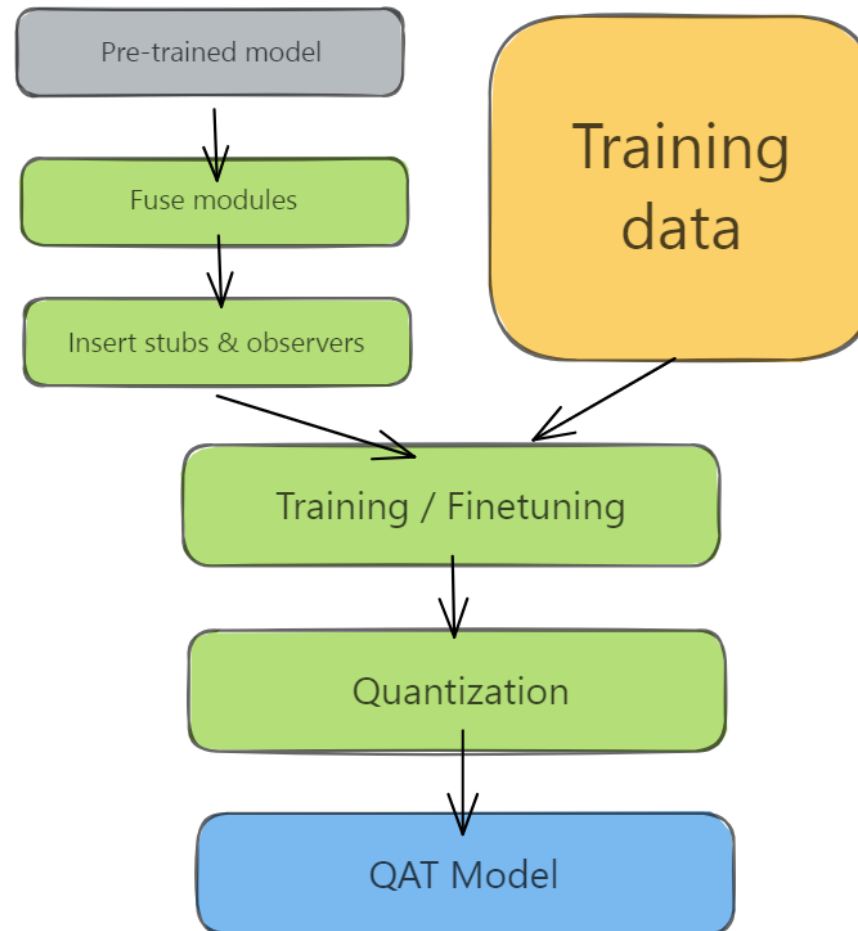
- $q = \text{clamp}[\text{round}(\frac{r}{s} + Z), -128, 127]$, where $s = \frac{\beta - \alpha}{\beta_q - \alpha_q}$ and $Z = \text{round}(\alpha_q - \frac{\alpha}{s})$
- For clip method, $[\alpha_q, \beta_q]$ is $[-256, 255]$ for int8 quantization.
- For weight, $[\alpha, \beta]$ is weight range.
- For activation, $[\alpha, \beta]$ is output range
- $r_q = (q - Z) * s$

```
FP32: tensor([ 1.4243, -0.7382, -1.0199,  0.6312, -0.4644, -1.2151,  1.1037,  0.2785])
INT8: tensor([ 0.7644, -0.5527, -0.5527,  0.6301, -0.4649, -0.5527,  0.7644,  0.2789])
INT8 integer representation: tensor([ 127., -128., -128.,  101., -111., -128.,  127.,  33.])
```

Introduction - PTQ




Introduction - QAT




Introduction - QAT

Bad performance



(a) Post training quantization

Good performance



(b) After quantization aware fine-tuning

Example code

```
✓ 2秒 [23] compare(model=FP32_model, device="cpu", test_loader=test_loader)
```

```
===== PERFORMANCE =====  
Size of the model(MB): 0.423146  
  
Accuracy: 8391/10000 (84%)
```

```
✓ 2秒 [24] compare(model=PTQ_model, device="cpu", test_loader=test_loader)
```

```
===== PERFORMANCE =====  
Size of the model(MB): 0.110646  
  
Accuracy: 8392/10000 (84%)
```

```
✓ 2秒 [25] compare(model=QAT_model, device="cpu", test_loader=test_loader)
```

```
===== PERFORMANCE =====  
Size of the model(MB): 0.110646  
  
Accuracy: 8469/10000 (85%)
```

Example code

Quantize layer by layer

```
===== PERFORMANCE =====  
Accuracy: 6655/10000 (67%)
```

Quantize at the same time

```
===== PERFORMANCE =====  
Accuracy: 8396/10000 (84%)
```

MSE

```
MSE of layer quantize_per_tensor is 0.5495123863220215  
MSE of layer nn1.relu is 1.3325119018554688  
MSE of layer nn2.relu is 1.7615503072738647  
MSE of layer dequantize is 14.916387557983398
```

Example code

```
def Calculate_scale_zero_point(x, mode="normal"):
    if mode == "normal":
        '''
        請完成以下程式碼
        '''

    elif mode == "clip":
        '''
        請完成以下程式碼
        '''

    return scale, zero_point
```

參考5、6頁的算法完成scale factor及zero point的計算

Example code

```
self.tensor = x
self.scale = scale
self.zero_point = zero_point
```

```
def _quantize(self, mode):
    if mode == "normal":
        self.qtensor_int = #請完成以下程式碼
        self.qtensor = #請完成以下程式碼

    elif mode == "clip":
        self.qtensor_int = #請完成以下程式碼
        self.qtensor_int = #請完成以下程式碼      clamp qtensor_int
        self.qtensor = #請完成以下程式碼
```

參考5、6頁的算法完成quantize的計算

Example code

```
[ ] scale_dic = []
    zero_dic = []

    #Calibrate to compute s、z of all layer at the same time
    for batch in train_loader:
        input, label = batch
        for node in ['x', 'relu', 'relu_1', 'nn3']:
            extractor = feature_extraction.create_feature_extractor(model, [node]).cpu()
            output = extractor(input)[node]
            q_min, q_max = -128, 127
            min_val, max_val = np.min(output.detach().numpy()), np.max(output.detach().numpy())
            scale = (max_val - min_val) / (q_max - q_min)
            zero = round(q_min - min_val / scale)
            q = Quantize_per_tensor(output, scale=scale, zero_point=zero, mode="normal")
            scale_dic.append(scale)
            zero_dic.append(zero)
        break

    print(scale_dic)
    print(zero_dic)
```

與Example當中quantize at the same time的作法一樣，先計算每個layer的scale、zero point

Example code

```
[ ] #define evaluate function
def Evaluate(model, loader):
    total = 0
    correct = 0
    with torch.no_grad():
        for data in loader:
            images, labels = data
            outputs = model(images)
            # the class with the highest energy is what we choose as prediction
            _, predicted = torch.max(outputs.data, 1)
            total += labels.size(0)
            correct += (predicted == labels).sum().item()

    test_loss = 0

    print("===== PERFORMANCE =====")
    print('\nAccuracy: {} / {} ({:.0f}%) \n'.format(correct, total, 100. * correct / total))

[ ] #Normal quantize
Evaluate(Quantized_normal_model, test_loader)

[ ] #Clip quantize
Evaluate(Quantized_clip_model, test_loader)
```

將兩種quantize計算方式的結果print出來並且放在結報中

Task

- Example code 30% (no need to write code)
 - Show accuracy comparison of FP32 model, PTQ model and QAT model
 - Self quantization
 - Show the MSE for each layer using two methods: quantizing layer by layer and quantizing all layers at the same time.
 - Show the difference of output distribution of above two method.
- Practice to implement quantization function 40% (Page 5)
 - Normal quantization
 - Clip quantization ($[\alpha_q, \beta_q]$ is $[-256, 255]$ but we only use $[-128, 127]$)
- Report and Question 30%

Task – cont.

- Report and Question 30%
 - Explain why the performance of quantize layer by layer is worse by using MSE and output distribution in the example code.
 - Besides the methods of scale factor and zero point mentioned in this lab, please provide the other example to determine the quantize value.